

1.1 Overall Objectives

In these laboratory sessions you will explore some of the computational intelligence techniques that have been covered in the first part of this course. These laboratories are self-explanatory and should serve as an introduction to the tools that will be required in your coursework. In particular, this laboratory series will enable you to:

- Write, run, and debug Python using Visual Studio
- Build a simple perception model of a neuron
- Assemble an artificial neural network using the perception model
- Implement back-propagation training for your artificial neural network
- Implement PCA and a k -nearest neighbour classification system

Important Message

This laboratory series builds on the knowledge developed within the lectures; you should make sure you are comfortable with the lecture content before attempting these laboratories. It does not assume any prior knowledge of Python and has a gentle introduction.

2.1 Week One Objectives

In this session you will learn how to create, run, and debug a Python program from within Microsoft Visual Studio. Your objectives are:

1. Create a Python project using either Visual Studio Code or Visual Studio
2. Run a basic Hello World example and become familiar with the editor
3. Write and test some simple programs that explore the key language concepts of Python
4. Learn how to use the debugger as a tool for software development

At the end of this laboratory session you should feel comfortable with creating, editing, debugging and running your own Python program. Note that this course has students from many different backgrounds and so while some will already be familiar with Python there are others who will never have used it before, and so this lab starts with the very basics.

2.2 Introduction

Personal Computers

Python and Visual Studio Code are both free software, so you can install them on your own computers for free. Python is available for pretty much every operating system and platform, but there are two caveats to consider:

1. Python underwent a major version shift between Python 2.7 and Python 3.x, the differences were big enough that Python 3.x is backwards incompatible with Python 2.7. In this course we will be using **Python 3.8** and this is the version that you should install on your machine.
2. Python itself is just the core language, and throughout this course we will make use of many different libraries that extend the core functionality. You can install different libraries individually, but the easiest way is to install a packaged distribution of Python and the major libraries, I recommend using Anaconda.

Thus, the best way to install Python on your home computer is to install Anaconda for Python 3.8 from: <https://www.anaconda.com/download/>

You are free to use whatever IDE (Integrated Development Environment) you wish. The University computers use Visual Studio and I recommend the free version called Visual Studio Code, available from: <https://code.visualstudio.com/>

Visual Studio Code and Anaconda support Windows, MacOS, and Linux. Basic instructions are provided via a short video on Moodle for running and debugging code, but the staff and lab demonstrators have experience of other platforms and will be happy to help. You can access the video introduction from:

<https://uniofbath.cloud.panopto.eu/Panopto/Pages/Viewer.aspx?id=6bcf89a9-6f64-4890-b7fc-ac4300e159a9>

University Computers

Python and Visual Studio are also available on the University computers, including in the library and via Remote Desktop to UniDesk. More information on connecting to UniDesk can be found from: <https://www.bath.ac.uk/guides/work-remotely-with-uniapps-and-unidesk/>

2.3 Some Examples

In order to introduce you to some of the key Python concepts I have provided some short code examples that show you the core aspects of the language. These are very simple, you do not need to try each one out, but you may find it useful to experiment with a few that you find more difficult. These examples can be used as a “cheat sheet” for more complicated programs. If you want to try them, you may either create new files for each example or you could add them sequentially to your existing .py file. Remember that you can print a variable to the console by using the *print* command, if you want to see what the contents might be.

Most of these examples are taken from the very good Python tutorial, you can look there for more detailed notes: <https://docs.python.org/3/tutorial/>

Variables and Mathematics

The first example to consider is how we store and manipulate variables (numbers, strings, etc). Python is very clever, and so unlike languages such as C it can work out what the type of your variable is from the contents – there is no need to explicitly type your variables. The code listing below shows some examples of creating and modifying some simple numerical variables.



Numerical

```
1  # This is a comment!
2
3  # Python knows that pi is a float
4  pi = 3.141
5
6  # and that freq is an integer
7  freq = 20000
8
9  # So nextFreq will also be an integer
10 nextFreq = freq + 1
11
12 # and radians/s will be a float, because we used a float in the calculation
13 rads = 2 * pi * nextFreq
14
15 # Classic division always returns a float
16 subRads = rads / 2.4
17
18 # Floor division returns the integer part (discards the fraction)
19 # Modulo division returns the fractional part (discards the integer part)
20 subRadsInt = rads // 2.4
21 remainder = rads % 2.4
22
23 # Powers can be calculated using **, here we square the frequency
24 f2 = freq ** 2
```

Code 2.1: numerical.py

Strings

```
1 # Strings can use either single or double quotes
2 firstName = "Bob"
3 lastName = 'Widlar'
4
5 # but be careful if you need to "contain" a quote mark, use \ to escape # lastName = Mc'Gee
6 lastName = 'Mc\Gee'
7
8 # String can be concatenated using +
9 fullName = firstName + " " + lastName
10
11 # Strings can be indexed, with the first character having index [0].
12 # There is no character type, a character is simply a string of length 1
13 firstLetter = fullName[0]
14
15 # You can also index from the right, so [-1] is the last character (-0 = 0)
16 lastLetter = fullName[-1]
17
18 # If you just want a slice, you can give a range # Starting at 0 and ending at 2
19 firstName = fullName[0:3]
20
21 # You can't change modify a string, but you can over-write it or make a new one # You can use len() to check the size..
22 nameLength = len(fullName)
```

Code 2.2: strings.py

Lists

```
1 # Lists are a way of storing lots of elements
2 squares = [1, 4, 9, 16, 25]
3
4 # Just like strings they can be indexed (from 0) and sliced
5 one = squares[0]
6 end = squares[-1]
7
8 # Note that the slice operation returns a new list containing the elements
9 squaresSlice = squares[0:3]
10
11 # Lists can be concatenated
12 squares = squares + [36, 49, 64, 81, 100]
13
14 # Unlike strings, lists are mutable. You can change the individual elements
15 squares[0] = 1 ** 2
16
17 # Len() also works on lists...
18 squaresLength = len(squares)
```

Code 2.4: lists.py

Loops and Flow Control

Loops and flow control are the next important area. You will have noticed by now that unlike C, Python does not require a semi-colon at the end of each line of code, or a set of curly brackets around code blocks. When want to group together bits of code, for example inside a loop, then we do so by indenting the code use a TAB space.

Loops

```
1  # Loops are a great way of repeating sections of code # Can you work out what this does?
2  a = 0
3  b = 1
4
5  # Importantly, the contents of the loop is enclosed using indents (single tab)
6  # Not {} like you get in C, loops can be nested, you just keep indenting
7  while b < 10:
8      print(b)
9      a = b
10     b = b + a
11
12 # We can also use for loops
13 # We need to generate a list of items to loop through
14 words = ['cat', 'window', 'defenestrate']
15 for word in words:
16     print(word)
17
18 # If we just want to loop a set number of times, we can make a range of integers # An easy way to this is using range()
19 # range(10) produces [0, 1, 2, 3, 4, 5, 6, 7, 8, 9]
20 for i in range(10):
21     print(i)
22
23 # We can perform a conditional check using the if statement
24 # the else and the elif (else if) are optional, note the indentation to group the code
25 if pi < 3:
26     pi = 3.141
27     print("Who ate the pie?")
28 elif pi > 4:
29     print("Too much pie..")
30 else:
31     pi = 0;
```

Code 2.5: loops.py

User Input

User Input

```
1  # We have already seen how to use print() to display data
2  # We can use input to get some data from the console
3  response = input("Is Python amazing? ")
4  print("You said: " + response)
5  # If we want to convert from string to int, or vice versa
6  limit = int(input("Enter a limit: "))
7  if limit > 3:
8      # We could also have used "Your limit of", limit, "was too big"
9      print("Your limit of " + str(limit) + " was too big")
```

Code 2.6: input.py

2.4 Debugging

One of the great features of Visual Studio Code is the excellent debugger, it will let us run our code one line at a time and inspect the value of all of our variables. In order to run the code line by line we need to insert a *breakpoint*, this tells Python to run from the start of our code until this line and stop there. You can insert a breakpoint by double clicking on the bar at the left-hand side of the editor window. Once in place a red circle will indicate the breakpoint, and when you run your code it will now stop at this point.

If you now hover the cursor over a variable it will show you the current value of that variable, you can also use the watch window (select Locals) to see the current value of all variables. Using the run options tool bar, you can either run to the next breakpoint, or single step the code one line at a time and watch as the variables change state.

2.5 Exercises

Spend some time familiarising yourself with the Visual Studio Code environment and the debugger, it will prove a valuable asset in the rest of the course. If you want some more challenging exercises, see if you can produce software to perform the following functions:

1. Given a number entered by the user, compute and display the factorial.
2. Store a secret password, ask the user to enter a password and compare the two (hint: strings are compared just like integers in C or MATLAB).
3. Display the maximum value that can be stored in a numerical variable.

1.1. Objectives

In this laboratory session you will learn how to create a Perceptron model of the neuron, generate some test vectors, and evaluate the performance of the Perceptron as a classification tool. Your objectives are:

1. Learn about the use of functions and libraries in Python
2. Create a function that is based on the Perceptron model of the neuron
3. Evaluate this function as a binary (Boolean) classification tool
4. Learn how to perform basic plotting in Python to graphically visualise data

At the end of this laboratory session you should have the basic framework for a Perceptron model, you are free to modify this in any way you want for your future coursework.

1.2. Introduction

We will be the same tools as the previous laboratory (i.e Visual Studio Code & the Anaconda Python distribution). Follow the instructions from Laboratory 1 to create a new project, be sure to select Anaconda as the Python environment to run.

The Perceptron

The Perceptron was first described in 1957 by Frank Rosenblatt at Cornell. It is a simple type of linear classifier - the output depends on a linear combination of the inputs. When first introduced, many Perceptrons were connected to form the first ever artificial neural network. In this laboratory we will start by considering a single Perceptron, whose output can be described by Equation 3.1:

$$f(x) = \begin{cases} 1, & \text{if } w \cdot x + b > 0 \\ 0, & \text{otherwise} \end{cases} \quad (3.1)$$

Where w is a vector of real-valued weights, $w \cdot x$ is the dot product and b is a bias term that can be used to scale the effective threshold. We can begin a simple analysis by constraining the inputs to be Boolean (0 or 1), and we can set the weights and the bias by hand. A graphical description of the model is given in Figure 3.1 for N inputs, and a mathematical expansion of Equation 3.1 is given in Equation 3.2, where θ represents the threshold function.

$$f(x) = \theta \left(\sum_{i=1}^N (w_i x_i + b) \right) \quad (3.2)$$

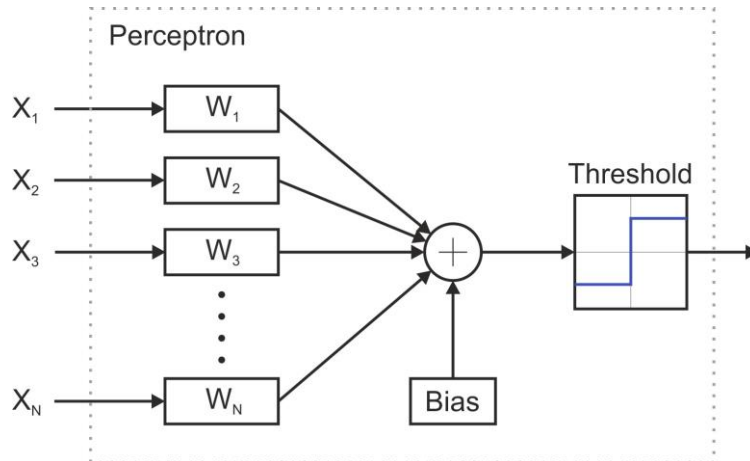


Figure 3.1: A single Perceptron with N inputs X_N , the inputs are multiplied by weights W_N and summed with the bias before a simple hardthreshold.

Python Libraries

In order to extend the capabilities of Python we are going to use some of the libraries that are provided by the Anaconda distribution (these libraries are actually written by the open source community, Anaconda is just a package that installs the most popular libraries. There are two libraries that we will make use of in this laboratory, NumPy and matplotlib.

Important Message

On the laboratory computers I have noticed that the first time you run a program that imports a library there can be quite a long delay before the program starts (around 6 seconds). This problem seems to be less noticeable on successive runs.

NumPy is a library that adds support for large, multi-dimensional arrays and matrices, along with a large collection of high-level mathematical functions to operate on these matrices. In many ways, NumPy adds a matrix feature set that resembles the capabilities of MATLAB, for example the dot product in Equation 3.1 can be performed in a single line of code, rather than writing a long for loop. This makes working with arrays and matrices much easier.

matplotlib is a plotting library that provides a good interface for drawing graphs and charts. As the name suggests it has been influenced by the plotting engine in MATLAB, as so the syntax is broadly similar and should be familiar to most of you.


You do not need to understand the intricacies of these libraries, fully worked examples will be given, however for your coursework you may choose to explore some of the capabilities in more detail.

1.3. Getting Started

If you have not already done so, create a new Python project using the instructions given in the lab from week 1.

Creating the Perceptron

Code listing 3.1 shows a simple program for simulating the Perceptron. Read through this code and type it into your own program. Library imports occur at the start of the program, followed by function definitions and then the main code. Note that in Python, functions work just like they do in MATLAB, here our Perceptron function takes three inputs and returns a single output. In the main code section, near the bottom, we create two lists that represent our inputs X_N and our weights W_N . Within the Perceptron function, we convert those lists into NumPy arrays. The reason for this is shown in line 14, using the NumPy array type we can compute the dot product in a single line, rather than using a for loop.

```
 Perceptron
1  # Import the NumPy library for matrix math
2  import numpy
3
4
5  # A single perceptron function
6  def perceptron(inputs_list, weights_list, bias):
7      # Convert the inputs list into a numpy array
8      inputs = numpy.array(inputs_list)
9
10     # Convert the weights list into a numpy array
11     weights = numpy.array(weights_list)
12
13     # Calculate the dot product
14     summed = numpy.dot(inputs, weights)
15
16     # Add in the bias
17     summed = summed + bias
18
19     # Calculate output
20     # N.B this is a ternary operator, neat huh?
21     output = 1 if summed > 0 else 0
22
23     return output
24
25  # Our main code starts here
26
27  # Test the perceptron
28  inputs = [1.0, 0.0]
29  weights = [1.0, 1.0]
30  bias = -1
31
32  print("Inputs: ", inputs)
33  print("Weights: ", weights)
34  print("Bias: ", bias)
35  print("Result: ", perceptron(inputs, weights, bias))
```

Code 3.1: perceptron.py

1.4. Exercises

Spend some time familiarising yourself with the basic Perceptron model. Once you have the model running, you should try the following exercises:

1. For a two-input perceptron with $W = [1.0, 1.0]$ and $b = -1$, simulate the output for all possible Boolean input combinations where $X \in \{0,1\}$.
2. Interpreting this information as a truth table, what logical function is being performed?
3. What are the weight vectors and the bias for the logical functions: AND, OR, NAND, NOR and XOR?
4. Why was 3. a trick question?

Further Exercises

Once you are happy with the basic Perceptron model, you can extend it using the matplotlib library to give you a graphical interpretation of the data. Code listing 3.2 shows a basic program that plots a scatter diagram of the input vector state space (every possible value of every input) for our two input Boolean Perceptron. Here are a few features you could add:

1. Modify your Perceptron code so that it automatically tests all possible input vectors (i.e $[0, 0]$, $[0, 1]$,...) *Hint: store the input vectors in a list, and use a for loop to test each input vector.*
2. With the aid of the example plotting code in listing 3.2, modify your code so that it automatically generates a plot of the input vector state space - use colour coding to signify if the Perceptron fired (output = 1) for that input.
3. Calculate the linear separator for the Perceptron and add it to the plot, an example of the final result is shown in Figure 3.2.




Figure 3.2: Example state space diagram with overlaid linear separator.

Plotting

```
1 # Import the matplotlib pyplot library
2 # It has a very long name, so import it as the name plt
3 import matplotlib.pyplot as plt
4
5 # Make a new plot (XKCD style)
6 fig = plt.xkcd()
7
8 # Add points as scatters - scatter(x, y, size, color)
9 # zorder determines the drawing order, set to 3 to make the
10 # grid lines appear behind the scatter points
11 plt.scatter(0, 0, s=50, color="red", zorder=3)
12 plt.scatter(0, 1, s=50, color="red", zorder=3)
13 plt.scatter(1, 0, s=50, color="red", zorder=3)
14 plt.scatter(1, 1, s=50, color="green", zorder=3)
15
16 # Set the axis limits
17 plt.xlim(-2, 2)
18 plt.ylim(-2, 2)
19
20 # Label the plot
21 plt.xlabel("Input 1")
22 plt.ylabel("Input 2")
23 plt.title("State Space of Input Vector")
24
25 # Turn on grid lines
26 plt.grid(True, linewidth=1, linestyle=':')
27
28 # Autosize (stops the labels getting cut off)
29 plt.tight_layout()
30
31 # Show the plot
32 plt.show()
```

Code 3.2: plotting.py

Week 3

The MLP

1.1. Objectives

In this laboratory session you will create a general purpose two-layer feed-forward artificial neural network. This neural network can solve many different classification problems, and it will extend the basic Perceptron from the previous laboratory. Your objectives are:

1. Learn about the use of classes and objects in Python
2. Create a neural network class that permits for an arbitrary number of neurons
3. Train the network using backpropagation
4. Explore the effect of the network size on training and performance

At the end of this laboratory session you should have a two-layer neural network, i.e a network with two layers of neurons, you are free to modify or extend this in any way you want for your future coursework.

1.2. Introduction

We will be using the same tools as the previous laboratory (i.e Visual Studio Code & the Anaconda Python distribution). Follow the instructions from Laboratory 1 to create a new project, be sure to select Anaconda as the Python environment to run.

Classes and Objects

One of the more powerful aspects of modern languages, such as Python, is the concept of classes and objects (these are also found in C++, which extended the C programming language). The idea behind an object is to group together data and functions that pertain to one "thing". The easiest way to understand objects is to see them in action, consider the code in listing 4.1. In this example we define a class (a class is the template that objects are based on). Our class defines a dog object, and that dog object has several functions. There is a special function called `__init__`, this is the function that is run when we make a new dog, it creates some variables that are specific to that individual dog object (using the *self* keyword). The class also defines some functions that the dog can have, these include a function called `bark`, which is arguably the most vital of all functions, if you are a dog.

The real power behind classes is the ability to create objects, every time we make a new object (a Dog), we are creating a copy of all the functions and variables that are defined by the class. Imagine if the class contained a neural network rather than a dog, then we would be able to easily create multiple neural networks and run them all at the same time, even though we only wrote the neural network code once. Figure 4.1 shows an example of what happens when we create a new object based on the class, you can think of the objects as *instances* and the classes as *templates*.




Figure 4.1: Example of two Dog objects being created based on the Dog class, which can be thought of as a template.

Dog Class

```

1  # Class template for a dog object
2  class Dog:
3      # Initialisation method, gets run whenever we create a new Dog
4      # The self just allows this function to reference variables relevant to this particular Dog
5      def __init__(self, name, hungerLevel):
6          self.name = name
7          self.hungerLevel = hungerLevel
8
9      # Query the status of the Dog
10     def status(self):
11         print("Dog is called ", self.name)
12         print("Dog hunger level is ", self.hungerLevel)
13         pass
14
15     # Set the hunger level of the dog
16     def setHungerLevel(self, hungerLevel):
17         self.hungerLevel = hungerLevel
18         pass
19
20     # Dogs can bark
21     def bark(self):
22         print("Woof!")
23         pass
24
25     # Create two dog objects
26     # Note that we don't need to include the self from the parameter list
27     lassie = Dog("Lassie", "Mild")
28     yoda = Dog("Yoda", "Ravenous")
29
30     # Check on Yoda & Lassie
31     yoda.status()
32     lassie.status()
33
34     # Get Lassie to bark
35     lassie.bark()
36
37     # Feed Yoda
38     yoda.setHungerLevel("Full")
39     yoda.status()

```

Code 4.1: dogClass.py

1.3. Getting Started

If you have not already done so, create a new Python project using the instructions given in the lab from week 1. Type in the example class code that defines the Dog class, and make sure you can run the code. Read through the comments and make sure you understand the principles at work before you continue with the neural network example.

Creating the Neural Network Class

In code listing 4.2 I have provided you with a complete neural network class that implements a two-layer feed-forward perceptron neural network. The number of inputs, the number of neurons, and the number of outputs is all definable when you create an instance (an object) of the neural network. For example, a neural network object with two inputs, two hidden neurons and one output neuron would look like that of Figure 4.2. This is approximately the configuration we decided was required (as a minimum) to solve the XOR problem.




Figure 4.2: Example two-layer feed-forward neural network with sigmoid activation functions.

The neural network has three functions, which are documented in Table 4.1. The basic pattern of use should be to create an instance of the network, to train the network using an inputs list and a target list, and to then query the network with some test data to validate the training. The default activation function is the logistic sigmoid (where $\theta \in [0,1]$) and there are no biases to the neurons.

Table 4.1: Functions of the NeuralNetwork Class

Function	Role
<code>N = NeuralNetwork(self, input_nodes, hidden_nodes, output_nodes, learning_rate)</code>	Creates a new neural network object with the number of inputs, hidden neurons and output neurons specified. An initial learning rate is also given and should be in the range $[0, 1]$.
<code>Train(self, inputs_list, targets_list)</code>	Performs one iteration of backpropagation training given a list of inputs and a list of targets. Note this does not work on lists of lists.
<code>Out = Query(inputs_list)</code>	Run the network with the given inputs.

ANN

```
1 # Import scipy.special for the sigmoid function expit()
2 import scipy.special, numpy
3
4 # Neural network class definition
5 class NeuralNetwork:
6     # Init the network, this gets run whenever we make a new instance of this class
7     def __init__(self, input_nodes, hidden_nodes, output_nodes, learning_rate):
8         # Set the number of nodes in each input, hidden and output layer
9         self.i_nodes = input_nodes
10        self.h_nodes = hidden_nodes
11        self.o_nodes = output_nodes
12
13        # Weight matrices, wih (input -> hidden) and who (hidden -> output)
14        self.wih = numpy.random.normal(0.0, pow(self.h_nodes, -0.5), (self.h_nodes, self.i_nodes))
15        self.who = numpy.random.normal(0.0, pow(self.o_nodes, -0.5), (self.o_nodes, self.h_nodes))
16
17        # Set the learning rate
18        self.lr = learning_rate
19
20        # Set the activation function, the logistic sigmoid
21        self.activation_function = lambda x: scipy.special.expit(x)
22
23    # Train the network using back-propagation of errors
24    def train(self, inputs_list, targets_list):
25        # Convert inputs into 2D arrays
26        inputs_array = numpy.array(inputs_list, ndmin=2).T
27        targets_array = numpy.array(targets_list, ndmin=2).T
28
29        # Calculate signals into hidden layer
30        hidden_inputs = numpy.dot(self.wih, inputs_array)
31
32        # Calculate the signals emerging from hidden layer
33        hidden_outputs = self.activation_function(hidden_inputs)
34
35        # Calculate signals into final output layer
36        final_inputs = numpy.dot(self.who, hidden_outputs)
37
38        # Calculate the signals emerging from final output layer
39        final_outputs = self.activation_function(final_inputs)
40
41        # Current error is (target - actual)
42        output_errors = targets_array - final_outputs
43
44        # Hidden layer errors are the output errors, split by the weights, recombined at hidden nodes
45        hidden_errors = numpy.dot(self.who.T, output_errors)
46
47        # Update the weights for the links between the hidden and output layers
48        self.who += self.lr * numpy.dot((output_errors * final_outputs * (1.0 - final_outputs)),
49        numpy.transpose(hidden_outputs))
50        # Update the weights for the links between the input and hidden layers
51        self.wih += self.lr * numpy.dot((hidden_errors * hidden_outputs * (1.0 - hidden_outputs)),
52        numpy.transpose(inputs_array))
53
54    # Query the network
55    def query(self, inputs_list):
56        # Convert the inputs list into a 2D array
57        inputs_array = numpy.array(inputs_list, ndmin=2).T
58
59        # Calculate signals into hidden layer
60        hidden_inputs = numpy.dot(self.wih, inputs_array)
61
62        # Calculate output from the hidden layer
63        hidden_outputs = self.activation_function(hidden_inputs)
64
65        # Calculate signals into final layer
66        final_inputs = numpy.dot(self.who, hidden_outputs)
67
68        # Calculate outputs from the final layer
69        final_outputs = self.activation_function(final_inputs)
70
71        return final_outputs
```

Code 4.2: ANN.py

1.4. Exercises

Spend some time familiarising yourself with the neural network model. You do not need to fully understand the internal workings, but you should try and understand the role of the three functions. Here are some exercises to try:

1. Create a neural network instance with two inputs, two hidden neurons and one output neuron. Create some test inputs and query the network (without any training), how do you explain the networks output?
2. Create a set of training vectors for all possible inputs, i.e $[[0,0], [0,1], [1,0], [1,1]]$, and a target vector for one of the logical functions such as AND.
3. Train the network once for each target, then query the network for all possible inputs. Has your network successfully learned the AND function? If not, then why not?
4. Extend your training routine to repeat the training process more than once for each target (using a loop), how many iterations does it take for the network to learn the AND function with two inputs?
5. Experiment with the number of neurons in the hidden layer, and the learning rate, what effect do these have on training the AND function?
6. If you have time, try other functions, such as XOR. What effect do the number of neurons in the hidden layer have and why?
7. How do these effects scale if you create, for example, a 4 input logical function?

1.1. Objectives

In this laboratory session you will use the general purpose two-layer MLP from laboratory two to perform automatic recognition of human handwriting. Specifically, you will devise and train an MLP to recognise digits from the MNIST database. Your objectives are:

1. Learn about the MNIST database and the format of the images
2. Learn how to read files in Python and display the MNIST images
3. Implement a two-layer MLP to perform classification on the dataset
4. Train the network, and evaluate the performance against a test set
5. Tweak the network parameters to try and improve the performance

At the end of this laboratory session you should have a practical neural network that can solve quite a complex task, this is the starting point for practical pattern recognition. You are free to modify or extend this in any way you want for your future coursework.

1.2. Introduction

We will be using the same tools as the previous laboratory (i.e Visual Studio Code & the Anaconda Python distribution). Follow the instructions from Laboratory 1 to create a new project, be sure to select Anaconda as the Python environment to run.

MNIST

The MNIST database (Modified National Institute of Standards and Technology database) is a large database of handwritten digits that is commonly used for training various image processing systems. The database is also widely used for training and testing in the field of machine learning. Recognising human handwriting is an ideal test for neural networks, because the problem is well defined but relatively difficult. The MNIST database contains 60,000 training images and 10,000 testing images. There have been a number of scientific papers on attempts to achieve the lowest error rate; one paper, using a hierarchical system of convolutional neural networks, manages to get an error rate on the MNIST database of 0.23%.



MNIST Files

The MNIST database in CSV format is available to download from Moodle. There are four files, the full test and training sets, and the smaller subsets of 100 training and 10 test sets.

Reading Files

The images are stored in Comma Separated Value (CSV) file format, which is easily readable in Python. You do not need to worry too much about how the data is represented, essentially the first value is the label (the actual number represented) and the subsequent values are a 28 x 28 array of pixels (giving a total of 784). To begin with, you should work on the smaller subset of MNIST, which provides you with 100 training examples and 10 test examples. In code listing 5.1, I have given you an example of how to open the smaller training set, read the first image, and display it using matplotlib. Figure 5.1 shows the resulting plot, in this case the first sample in the training set is the number 5. You should run this example and take a look at the first few images from both of the smaller training and test sets.

Draw Number

```
1 # Import numpy for arrays and matplotlib for drawing the numbers
2
3 import numpy
4 import matplotlib.pyplot as plt
5
6 # Open the 100 training samples in read mode
7 data_file = open("mnist_train_100.csv", 'r')
8
9 # Read all of the lines from the file into memory
10 data_list = data_file.readlines()
11
12 # Close the file (we are done with it)
13 data_file.close()
14
15 # Take the first line (data_list index 0, the first sample), and split it up based on the commas
16 # all_values now contains a list of [label, pixel 1, pixel 2, pixel 3, ... ,pixel 784]
17 all_values = data_list[0].split(',')
18
19 # Take the long list of pixels (but not the label), and reshape them to a 2D array of pixels
20 image_array = numpy.asarray(all_values[1:]).reshape((28, 28))
21
22 # Plot this 2D array as an image, use the grey colour map and don't interpolate
23 plt.imshow(image_array, cmap='Greys', interpolation='None')
24 plt.show()
```

Code 5.1: drawNumber.py




Figure 5.1: Greyscale plot of the first example from the training set of 100 numbers (in this case, the number 5).

Table 5.1: Suggested parameters for two-layer MNIST MLP.

Parameter	Suggested Value
Input Nodes	784
Hidden Nodes	100
Output Nodes	10
Learning Rate	0.3

1.3. Getting Started

If you have not already done so, create a new Python project using the instructions given in the lab from week 1. Type in the example code that plots the handwritten numbers, and make sure you can run the code. Read through the comments and make sure you understand the principles at work before you continue with the neural network example.

Configuring the Network

Before we can jump in and start using our neural network to classify the images, we need to think about the topology of network that we want to use. We will also need to think about the format of the input and output data. We know that the input will be a vector of 784 pixels, so it is easy to set the number of input nodes. The output needs to tell us what number the network thinks was in the image, the easiest way to do this is to have 10 output nodes, each one corresponding to the probability that the image contained that number. So, for example a value of 0.99 at output node 0 would mean that the network was certain that the image contained the number 0. The number of hidden nodes and the learning rate are, at this stage, best guess. I have put some starting values in Table 5.1.

Training and Testing the Network

In Code Listing 5.2, is the code to train the network using the 100 sample training set (note that in the code the neural network is called *n*). There are two important aspects to note. Firstly, the pixels in each image are in the range 0...255, I have rescaled them to be in the range 0.01..1. It is often helpful to have non-zero inputs, as a zero input can cause the training process to become stuck (as changing to associated weight has no effect on the output). The second key aspect is that I create a target vector with 0.01 indicating “not this digit” and 0.99 indicating “this digit”. In the previous labs you should have observed that the networks output never reaches 1 or 0 exactly, this is because the logistic sigmoid function is *asymptotic* - if we set targets of 1 or 0 then the training process becomes saturated at the asymptote. Therefore, you may have noticed that your MLP performed worse at the simple logic functions than you hand crafted Perceptron, which used the hard threshold function. In Code Listing 5.3, I have provided example code that tests the network using the 10-sample test set. A scorecard is created, with a 1 appended for each correct classification, and a 0 for an incorrect classification. Finally, the performance score is calculated as the percentage of correct classifications.

Training

```
1 # Load the MNIST 100 training samples CSV file into a list
2 training_data_file = open("mnist_train_100.csv", 'r')
3 training_data_list = training_data_file.readlines()
4 training_data_file.close()
5
6 # Train the neural network on each training sample
7 for record in training_data_list:
8     # Split the record by the commas
9     all_values = record.split(',')
10    # Scale and shift the inputs from 0..255 to 0.01..1
11    inputs = (numpy.asarray(all_values[1:]) / 255.0 * 0.99) + 0.01
12    # Create the target output values (all 0.01, except the desired label which is 0.99)
13    targets = numpy.zeros(output_nodes) + 0.01
14    # All_values[0] is the target label for this record
15    targets[int(all_values[0])] = 0.99
16    # Train the network
17    n.train(inputs, targets)
18 pass
```

Code 5.2: mnistTrain.py

Testing

```
1 # Load the MNIST test samples CSV file into a list
2 test_data_file = open("mnist_test_10.csv", 'r')
3 test_data_list = test_data_file.readlines()
4 test_data_file.close()
5
6 # Scorecard list for how well the network performs, initially empty
7 scorecard = []
8
9 # Loop through all of the records in the test data set
10 for record in test_data_list:
11     # Split the record by the commas
12     all_values = record.split(',')
13     # The correct label is the first value
14     correct_label = int(all_values[0])
15     print(correct_label, "Correct label")
16     # Scale and shift the inputs
17     inputs = (numpy.asarray(all_values[1:]) / 255.0 * 0.99) + 0.01
18     # Query the network
19     outputs = n.query(inputs)
20     # The index of the highest value output corresponds to the label
21     label = numpy.argmax(outputs)
22     print(label, "Network label")
23     # Append either a 1 or a 0 to the scorecard list
24     if (label == correct_label):
25         scorecard.append(1)
26     else:
27         scorecard.append(0)
28     pass
29 pass
30
31 # Calculate the performance score, the fraction of correct answers
32 scorecard_array = numpy.asarray(scorecard)
33 print("Performance = ", (scorecard_array.sum() / scorecard_array.size)*100, '%')
```

Code 5.3: mnistTest.py

1.4. Exercises

Spend some time familiarising yourself with the code for training and testing the network with the MNIST dataset. You do not need to fully understand the internal workings, but you should understand the basic flow of the code. Here are some exercises to try:

1. Using the template network class from week 3, create a neural network with the parameters given in Table 5.1 Note: you will need to create variables for the network parameters such as output nodes.
2. Train this network using the MNIST100 training set, and then test it using the MNIST10 testing set. What is your networks performance? Mine was 60%, but there will be small variances due to the random initial weights.
3. Display the images that were incorrectly classified, do you think that you (as a human), could have got those classifications correct?
4. Modify the network topology, the learning rate and any other parameters you desire – what is the highest network performance you can get?
5. Re-run the process using the full size MNIST training and testing sets. This may take some time. What is the performance now? Why is it different than before?

1.1. Objectives

In this laboratory session you will use Principal Component Analysis (PCA) and the k-nearest neighbour (KNN) algorithms to perform automatic recognition of human handwriting. Specifically, you will use PCA to automatically extract a set of features, and KNN to classify digits from the MNIST dataset using the features. Your objectives are:

1. Load the MNIST images, and learn about the CI tools within the scikit-learn package
2. Perform PCA on the training set and evaluate the variance distribution
3. Use the principal components from objective 2 to create a KNN classifier
4. Extract the same principal components from the test data set, and classify them using the KNN classifier
5. Evaluate the overall performance and speed, and experiment with the different parameters

At the end of this laboratory session you should have a working handwriting recognition system that you can directly compare to the neural network system in the previous labs. You will also have a working demonstration of both PCA and KNN. You are free to modify or extend this in any way you want for your future coursework.

1.2. Introduction

We will be the same tools as the previous laboratory (i.e Visual Studio Code & the Anaconda Python distribution). Follow the instructions from Laboratory 1 to create a new project, be sure to select Anaconda as the Python environment to run.

MNIST

We will be using the same MNIST database as Week 4, as a reminder MNIST is a large database of handwritten digits that is commonly used for training various image processing systems. The MNIST database contains 60,000 training images and 10,000 testing images. There have been many scientific papers on attempts to achieve the lowest error rate; one paper, using a hierarchical system of convolutional neural networks, manages to get an error rate on the MNIST database of 0.23%.



MNIST Files

The MNIST database in CSV format is available to download from Moodle. There are four files, the full test and training sets, and the smaller subsets of 100 training and 10 test sets.

Scikit-learn

In the laboratories so far, you have created most of the functional code yourselves, making use of libraries such as numpy and matplotlib to perform useful functions. This has been a worthwhile exercise because it enables a far greater level of understanding than simply using someone else's code. However, Computational Intelligence and machine learning are very popular areas of research and there are many Python libraries that provide you with working examples of structures such as neural networks. In this laboratory, we will be using the PCA algorithm and the KNN classifier. We could write the code for these from scratch, but it would become very long and quite tedious. Instead we will make use of a very popular Python library called scikit-learn. Scikit-learn is a machine learning library that includes various classification, regression and clustering algorithms including PCA and KNN. There is plenty of documentation available on the scikit-learn website if you wish to explore the algorithms further.

1.3. Getting Started

If you have not already done so, create a new Python project using the instructions given in the lab from week 1.

Principal Component Analysis

Our first step is to look at the code that reads in the MNIST data and performs PCA. In code listing [6.1](#) the two CSV files are loaded, making use of a numpy file read command. Next the data and labels are identified and a PCA object is initiated with a parameter that sets the number of components to extract. The training data is then fed into the PCA object, once the principal components have been extracted it is possible to plot the variance that is represented by the components (if the components represented all the variance, then the sum of the variances would be 1).

PCA

```
1  # Numpy for useful maths
2  import numpy
3  # Sklearn contains some useful CI tools
4  # PCA
5  from sklearn.decomposition import PCA
6  from sklearn.preprocessing import MinMaxScaler
7  # k Nearest Neighbour
8  from sklearn.neighbors import KNeighborsClassifier
9  # Matplotlib for plotting
10 import matplotlib.pyplot as plt
11
12 # Load the train and test MNIST data
13 train = numpy.loadtxt('mnist_train_100.csv', delimiter=',')
14 test = numpy.loadtxt('mnist_test_10.csv', delimiter=',')
15
16 # Separate labels from training data
17 train_data = train[:, 1:]
18 train_labels = train[:, 0]
19 test_data = test[:, 1:]
20 test_labels = test[:, 0]
21
22 # Select number of components to extract
23 pca = PCA(n_components = 10)
24 # Fit to the training data
25 pca.fit(train_data)
26
27 # Determine amount of variance explained by components
28 print("Total Variance Explained: ", numpy.sum(pca.explained_variance_ratio_))
29
30 # Plot the explained variance
31 plt.plot(pca.explained_variance_ratio_)
32 plt.title('Variance Explained by Extracted Components')
33 plt.ylabel('Variance')
34 plt.xlabel('Principal Components')
35 plt.show()
```

Code 6.1: PCA.py

k-Nearest Neighbour

Once we have setup our PCA, we can move on the classification stage. Code listing [6.2](#) shows the next steps, the training and test data are both transformed using PCA into a list of principal component scores, these scores are then normalised. Next, we create a KNN classifier with the training principal component scores and labels as the training data, note that the KNN is initiated with $k = 5$ and the $p = 2$ (Euclidean) norm. Finally, we can now use the KNN object to predict the classification for the test data. As with the previous laboratories a scorecard is created, and the percentage success rate is calculated.

KNN

```
1 # Extract the principal components from the training data
2 train_ext = pca.fit_transform(train_data)
3 # Transform the test data using the same components
4 test_ext = pca.transform(test_data)
5
6 # Normalise the data sets
7 min_max_scaler = MinMaxScaler()
8 train_norm = min_max_scaler.fit_transform(train_ext)
9 test_norm = min_max_scaler.fit_transform(test_ext)
10
11 # Create a KNN classification system with k = 5
12 # Uses the p2 (Euclidean) norm
13 knn = KNeighborsClassifier(n_neighbors=5, p=2)
14 knn.fit(train_norm, train_labels)
15
16 # Feed the test data in the classifier to get the predictions
17 pred = knn.predict(test_norm)
18
19 # Check how many were correct
20 scorecard = []
21
22 for i, sample in enumerate(test_data):
23     # Check if the KNN classification was correct
24     if round(pred[i]) == test_labels[i]:
25         scorecard.append(1)
26     else:
27         scorecard.append(0)
28     pass
29
30 # Calculate the performance score, the fraction of correct answers
31 scorecard_array = numpy.asarray(scorecard)
32 print("Performance = ", (scorecard_array.sum() / scorecard_array.size) * 100, ' % ')
```

Code 6.2: KNN.py

1.4. Exercises

Spend some time familiarising yourself with the code for extracting the principal component scores and the KNN classifier. You do not need to fully understand the internal workings, but you should understand the basic flow of the code. Here are some exercises to try:

1. Try increasing the number of principal components, how many are needed to explain all the variance within the training data?
2. How many principal components do you think is enough to perform effective classification?
3. What is the best performance you can get on the MNIST100/10 data set? What happens if you change the parameters k and p ?
4. How does your best performance compare to the neural network from laboratory four?
5. What can you say about the time taken to run this type of classifier?
6. Using the code you have been provided, use PCA to reduce the number of features and then use a neural network to classify the MNIST characters. Does this give you the best result?

Lab 6: Optimization using Simulated Annealing

1 Objectives

The main objective for this lab is to create a numerical optimization example using simulated annealing and then to develop your own custom version of the optimizer for more general purpose applications.

1. Learn how simulated annealing optimization works using an example
2. Create a simple simulated annealing function to optimize a number
3. Create a simulated annealing function to optimize a function

At the end of this laboratory session you should have the basic framework for simulated annealing optimization and you could explore how to potentially apply this to optimization of neural network parameters.

2 Getting Started

As we have seen in the lecture on optimization, the algorithm for simulated is as shown in figure 1. The key aspect of the algorithm is that we need to have specific functions to create a new (random) starting point, a *neighbour* version for comparison based on this starting individual, a *cost* function, an acceptance probability function, and finally the overall annealing loop.

The good news as far as this is concerned is that if we keep the functions generic, then all we need to do is to modify the cost function and creation of new values to optimize and the main loop function can remain unchanged.

In the lecture we saw a simple example where a 10x10 matrix of energy demand figures were used as the target for optimization of a 10x10 matrix of energy supply values, and this was based on the python code below:

Listing 1: sa.py

```
1 import numpy as np
2 import math
3 import random
4 import matplotlib.pyplot as plt
```

```

5
6 def acceptance_probability(old_cost , new_cost , T):
7     a = math.exp((old_cost-new_cost)/T)
8     return a
9
10 def anneal(solution , target , alpha , iterations , var):
11     old_cost = cost(solution , target)
12     cost_values = list()
13     cost_values.append(old_cost)
14     T = 1.0
15     T_min = 0.000001
16     while T > T_min:
17         i = 1
18         while i <= iterations:
19             print("Iteration_:_" + str(i) + "Cost_:_" + str(old_cost))
20             new_solution = neighbour(solution , var)
21             new_cost = cost(new_solution , target)
22             ap = acceptance_probability(old_cost , new_cost , T)
23             if ap > random.random():
24                 solution = new_solution
25                 old_cost = new_cost
26                 i += 1
27             cost_values.append(old_cost)
28             T = T*alpha
29     return solution , old_cost , cost_values
30
31 def cost(supply , demand):
32     delta = np.subtract(demand , supply)
33     delta2 = np.square(delta)
34     ave = np.average(delta2)
35     dcost = math.sqrt(ave)
36     return dcost
37
38 def neighbour(solution , d):
39     delta = np.random.random((10 , 10))
40     scale = np.full((10 , 10) , 2*d)
41
42     offset = np.full((10 , 10) , 1.0 - d)
43
44     var = np.multiply(delta , scale)
45     m = np.add(var , offset)
46
47     new_solution = np.multiply(solution , m)
48
49     return new_solution
50
51
52
53 supply = np.full((10 , 10) , 0.5)
54 demand = np.random.random((10 , 10))

```

```

55
56 print("Supply")
57 print(supply)
58
59 print("Demand")
60 print(demand)
61
62 rms = cost(supply,demand)
63 print("RMS=_"+ str(rms))
64
65 alpha=0.9
66 iterations = 200
67 var = 0.01
68
69 final_solution , cost , cost_values = anneal(supply,demand,alpha,iterations,var)
70
71 print(demand)
72 print(final_solution)
73 print(cost)
74
75 plt.subplot(232)
76 plt.title("Error_Function_in_Simulated_Annealing")
77 plt.plot(cost_values)
78 plt.grid(True)
79
80 plt.subplot(234)
81 plt.title("Initial_Supply_Matrix")
82 plt.imshow(supply, cmap='hot')
83
84 plt.subplot(235)
85 plt.title("Optimized_Supply_Matrix")
86 plt.imshow(final_solution, cmap='hot')
87
88 plt.subplot(236)
89 plt.title("Demand_Matrix")
90 plt.imshow(demand, cmap='hot')
91
92 plt.show()

```

Now, in this example we have quite a bit of plotting of data, but we can examine the code section by section to understand what is going on.

2.1 Simulated Annealing Function

The heart of the simulated annealing algorithm is implemented in the `anneal` function. the current solution is passed to the function, along with the target value (this would be a constraint perhaps or minima/maxima for example), `alpha` (reduction in temperature per iteration), `iterations` (number of iterations at each temperature) and `var` (the variation of the random number used to

generate a new neighbour).

The temperature is defined inside this routine and you could legitimately move this outside and make it a parameter.

It would also be good practice to consider making a simulated annealing class and have all the relevant functions defined in the class.

Listing 2: anneal function

```
1 def anneal(solution, target, alpha, iterations, var):
2     old_cost = cost(solution, target)
3     cost_values = list()
4     cost_values.append(old_cost)
5     T = 1.0
6     T_min = 0.000001
7     #alpha = 0.9
8     #iterations = 100
9     # var = 0.01
10    while T > T_min:
11        i = 1
12        while i <= iterations:
13            print("Iteration: " + str(i) + " Cost: " + str(old_cost))
14            new_solution = neighbour(solution, var)
15            new_cost = cost(new_solution, target)
16            ap = acceptance_probability(old_cost, new_cost, T)
17            if ap > random.random():
18                solution = new_solution
19                old_cost = new_cost
20            i += 1
21            cost_values.append(old_cost)
22        T = T*alpha
23    return solution, old_cost, cost_values
```

The cost function is defined as the difference in the value using the current solution and the desired target. In this example it is simply the RMS difference of all the values, and we are looking to minimize the cost as a result. In the annealing function we create a list called cost_values and the rationale for this is so we have a complete record of all the cost values which can be plotted at the conclusion of the optimization.

2.2 Acceptance Probability function

The acceptance probability function is the implementation of the exponential calculation of the probability that we will accept the new value, and this is useful as it can be applied to any problem, as long as we can reduce the cost to a single value. We could therefore use this in ANY simulated annealing implementation.

Listing 3: acceptance probability

```
1 def acceptance_probability(old_cost, new_cost, T):
2     a = math.exp((old_cost - new_cost)/T)
3     return a
```

2.3 Neighbour function

The neighbour function in this case is a new 10x10 matrix of values that is a random variation of the current solution. This function is specific to the current example and will need a new routine to be made to suit the problem under investigation.

Listing 4: new neighbour

```
1 def neighbour(solution,d):
2     delta = np.random.random((10,10))
3     scale = np.full((10,10),2*d)
4
5     offset = np.full((10,10),1.0-d)
6
7     var = np.multiply(delta,scale)
8     m = np.add(var,offset)
9
10    new_solution = np.multiply(solution,m)
11
12    return new_solution
```

2.4 Cost function

The cost function in this example is a calculation of the RMS error between the energy supply and demand matrices. This has specific names for the supply and demand (but easy to change to target and current or similar to make it general) and it is based on np matrices, therefore these would need to be the format in which the data sets were defined (but obviously this could be done for most situations).

The cost function used will depend on the application and how you wished to calculate it (for example, Euclidean distance might be used)

Listing 5: cost function

```
1 def cost(supply,demand):
2     delta = np.subtract(demand,supply)
3     delta2 = np.square(delta)
4     ave = np.average(delta2)
5     dcost = math.sqrt(ave)
6     return dcost
```

3 Lab Exercise

1. Modify the code (or create your own from scratch) to create a simple optimization where the function of the optimizer is to simply find a number specified in the shortest number of iterations

2. How do the number of iterations, variability and alpha parameters affect how quickly a solution is found?
3. Modify the code (or create your own from scratch) to create an optimization of a linear array of 10 numbers and find the values as quickly as possible
4. Create an optimization to optimize a straight line that gives the best fit of a random array of 100 points in XY space between 0 and 100 in both axes.

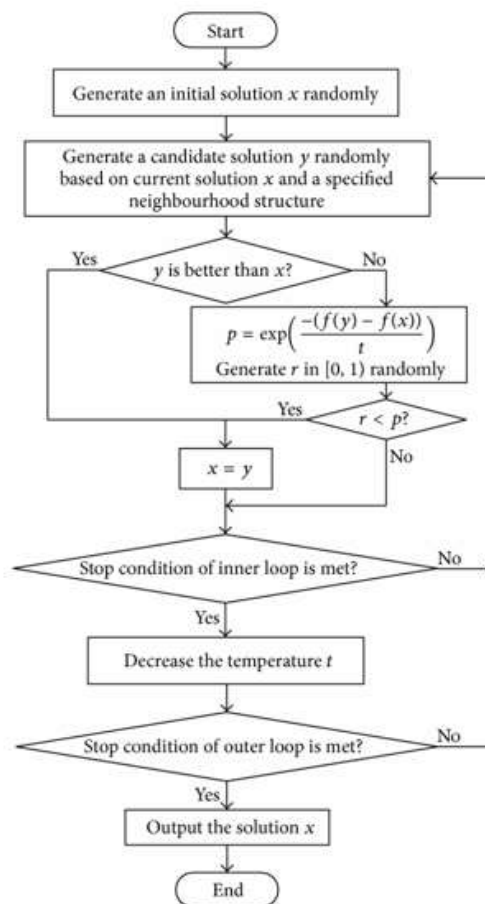


Figure 1: Algorithm for Simulated Annealing

Lab 8: Using Fuzzy Logic

1 Objectives

The main objective for this lab is to explore how Fuzzy Logic Systems work and experiment with the Fuzzy characteristics. We will use a standard example which controls the water level in a tank using a fuzzy valve controller.

In this lab you will have a break from Python and will use Matlab

1. Run Matlab
2. Load the Fuzzy Logic Water Tank Controller
3. Run simulations of the Fuzzy Logic Water Tank Controller
4. Experiment with the fuzzy logic parameters

At the end of this laboratory session you should have a better understanding of the nature of fuzzy logic systems and be able to see how they operate.

2 Getting Started

We will be using Matlab, not Python, in this lab, so run Matlab from one of the lab computers. In the Matlab command line please load up the fuzzy logic controller:

```
fuzzy('tank')
```

This will load up the fuzzy logic controller as shown in Figure 1.

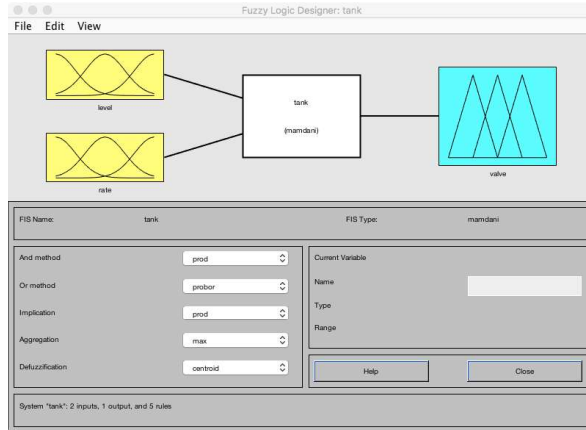


Figure 1: Fuzzy Controller for Water Tank Example

leave this for now and next load up the simulink model of the water tank system using the following command:

sltank

This will then load up the simulink model as shown in Figure ??.

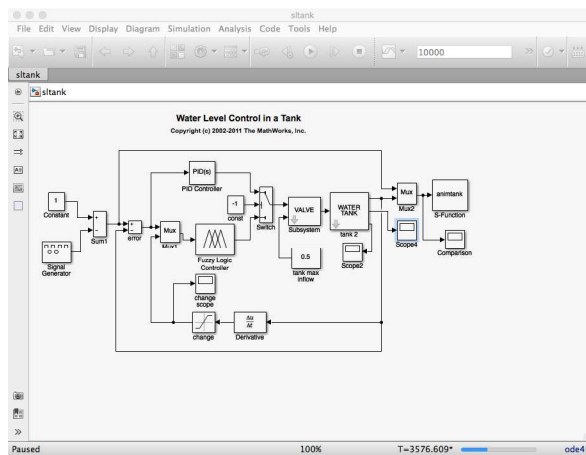


Figure 2: Water Tank Control System Example

You can see the fuzzy logic controller implemented as a block in the simulink design and the simulink model can be simulated to observe the behaviour of the controller using the fuzzy logic controller in Figure ??.




Figure 3: Water Tank Control System Example Water Level

The membership functions can also be plotted by double clicking on the fuzzy logic interface and they will popup in a new window.

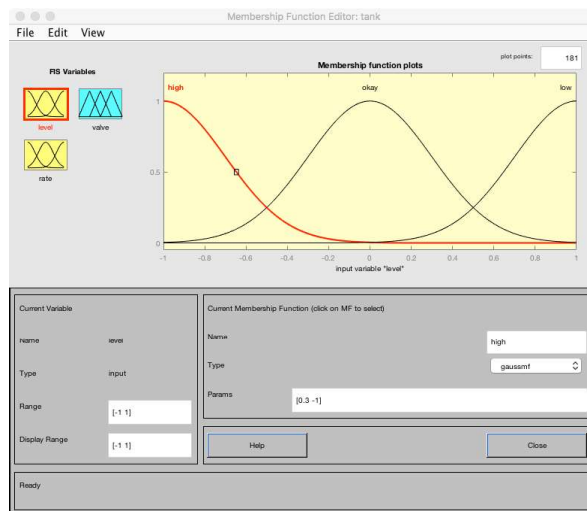


Figure 4: Fuzzy logic Membership functions

3 Comparison with a standard PID controller

The Simulink model allows the use of a PID controller, by changing the constant selector input to the switch model. If the switch is set to -1, then the Fuzzy controller is used, however +1 means the PID controller is used.

Try experimenting with the PID parameters and see the effect this has on the control.

4 Experimenting with the fuzzy logic controller

The fuzzy logic controller can be modified by clicking on the membership function settings and the values of each set can be modified as can the function (the default is Gaussian, but can be changed to trapezoidal or triangular for example).

Is there a difference between Gaussian, Triangular and Trapezoidal functions?

What happens if you change the values of the membership functions?